

Databázové systémy

Zadanie 6 – ORM

Meno: Andrii Rybak

AIS ID: 105840

Cvičiaci: Ing. Vladimír Kunštár

Akademický rok: 2021/2022

Content

1. /patches.....	3
Initial:.....	3
ORM generated:	3
Summary	4
2. /players/14944/game_exp.....	5
Initial:.....	5
ORM generated:	6
Summary	7
3. /players/14944/game_objectives	8
Initial:.....	8
ORM generated:	8
Summary	10
4. /players/14944/abilities	11
Initial:.....	11
ORM generated:	13
Summary	15
5. /matches/21421/top_purchases/	16
Initial:.....	16
ORM generated:	17
Summary	19
6. /abilities/5004/usage/	20
Initial:.....	20
ORM generated:	21
Summary	23
7. /statistics/tower_kills/	24
Initial:.....	24
ORM generated:	25
Summary	26

1. /patches

Initial:

```
WITH table1
  AS (SELECT NAME
        patch_version,
        Cast(Extract(epoch FROM release_date) AS INTEGER) AS
        patch_start_date
      FROM patches
      ORDER BY patch_version ASC),
  table2
  AS (SELECT patch_version,
        patch_start_date,
        Lead(patch_start_date, 1)
          OVER (
            ORDER BY patch_version) patch_end_date
      FROM table1)
SELECT patch_version,
       patch_start_date,
       patch_end_date,
       mts.id AS match_id,
       Round(mts.duration / 60.00, 2) AS match_duration
FROM   table2
      LEFT JOIN matches AS mts
        ON mts.start_time >= patch_start_date
        AND mts.start_time < patch_end_date
```

	QUERY PLAN
	text
1	Nested Loop Left Join (cost=1.69..18560.07 rows=105556 width=76) (actual time=1917.938..19475.041 rows=50005 loops=1)
2	[...] Join Filter: (((mts.start_time >= ((date_part('epoch'::text, patches.release_date)::integer)) AND (mts.start_time < (lead(((date_part('epoch'::text, patches.release_date)::integer), 1) OVER (????))))
3	[...] Rows Removed by Join Filter: 900000
4	[...] -> WindowAgg (cost=1.69..2.21 rows=19 width=40) (actual time=0.554..1.326 rows=19 loops=1)
5	[...] -> Sort (cost=1.69..1.74 rows=19 width=36) (actual time=0.506..0.717 rows=19 loops=1)
6	[...] Sort Key: patches.name
7	[...] Sort Method: quicksort Memory: 25kB
8	[...] -> Seq Scan on patches (cost=0.00..1.28 rows=19 width=36) (actual time=0.041..0.257 rows=19 loops=1)
9	[...] -> Materialize (cost=0.00..1266.00 rows=50000 width=12) (actual time=0.011..525.040 rows=50000 loops=19)
10	[...] -> Seq Scan on matches mts (cost=0.00..1016.00 rows=50000 width=12) (actual time=0.022..487.312 rows=50000 loops=1)
11	Planning Time: 0.180 ms
12	Execution Time: 19953.505 ms

ORM generated:

```
SELECT anon_1.patch_version AS anon_1_patch_version,
       anon_1.patch_start_date AS anon_1_patch_start_date,
       anon_1.patch_end_date AS anon_1_patch_end_date,
       matches.id AS match_id,
       Round(matches.duration / 60.0, 2) AS match_duration
FROM   (SELECT anon_2.patch_version AS patch_version,
               anon_2.patch_start_date AS patch_start_date,
               Lead(anon_2.patch_start_date, 1)
                 OVER (
                   ORDER BY anon_2.patch_version ) AS patch_end_date
      FROM   (SELECT patches.NAME
               AS
               patch_version,
               Cast(Extract(epoch FROM patches.release_date) AS INTEGER) AS
               patch_start_date
      FROM   patches) AS anon_2) AS anon_1
      LEFT OUTER JOIN matches
        ON matches.start_time >= anon_1.patch_start_date
        AND matches.start_time < anon_1.patch_end_date
```

QUERY PLAN		
text		
1	Nested Loop Left Join (cost=1.59..18559.98 rows=105556 width=76) (actual time=2063.780..19368.389 rows=50005 loops=1)	
2	[...] Join Filter: (((matches.start_time >= ((date_part('epoch':text, patches.release_date)::integer)) AND (matches.start_time < (lead((date_part('epoch':text, patches.release_date)::integer, 1) OVER (???))	
3	[...] Rows Removed by Join Filter: 900000	
4	[...] -> WindowAgg (cost=1.59..2.12 rows=19 width=40) (actual time=0.682..1.583 rows=19 loops=1)	
5	[...] -> Sort (cost=1.59..1.64 rows=19 width=40) (actual time=0.592..0.793 rows=19 loops=1)	
6	[...] Sort Key: patches.name	
7	[...] Sort Method: quicksort Memory: 25kB	
8	[...] -> Seq Scan on patches (cost=0.00..1.19 rows=19 width=40) (actual time=0.037..0.300 rows=19 loops=1)	
9	[...] -> Materialize (cost=0.00..1266.00 rows=50000 width=12) (actual time=0.011..520.889 rows=50000 loops=19)	
10	[...] -> Seq Scan on matches (cost=0.00..1016.00 rows=50000 width=12) (actual time=0.019..524.267 rows=50000 loops=1)	
11	Planning Time: 0.196 ms	
12	Execution Time: 19834.763 ms	

Summary

ORM generated almost the same query as original one, so it is impossible to get significantly different performance from it. The only difference is that in original query was used WITH to make subtable, but in the generated query was used nested SELECT. However, query plan has the same structure and almost the same values in both cases.

The most expensive operation is left join, which took ~19200 msec for all rows. Join operation took more than 90% of all query execution time.

2. /players/14944/game_exp

Initial:

```
SELECT pl.id,
       CASE
         WHEN pl.nick IS NULL THEN 'unknown'
         ELSE pl.nick
       end AS player_nick,
       hr.localized_name AS hero_localized_name,
       Round(mt.duration / 60.00, 2) AS match_duration_minutes,
       Coalesce(mt_pl_dt.xp_hero, 0)
+ Coalesce(mt_pl_dt.xp_creep, 0)
+ Coalesce(mt_pl_dt.xp_other, 0)
+ Coalesce(mt_pl_dt.xp_roshan, 0) AS experiences_gained,
       mt_pl_dt.level AS level_gained,
       CASE
         WHEN player_slot >= 0
           AND player_slot <= 4
           AND radiant_win IS TRUE THEN true
         WHEN player_slot >= 128
           AND player_slot <= 132
           AND radiant_win IS FALSE THEN true
         ELSE false
       end AS winner,
       mt.id AS match_id
FROM   players AS pl
       JOIN matches_players_details AS mt_pl_dt
         ON mt_pl_dt.player_id = pl.id
       JOIN heroes AS hr
         ON mt_pl_dt.hero_id = hr.id
       JOIN matches AS mt
         ON mt_pl_dt.match_id = mt.id
WHERE  pl.id = 14944
ORDER BY match_id ASC
```


	QUERY PLAN	
	text	
1	Gather Merge (cost=14228.08..14229.15 rows=9 width=91) (actual time=26.443..29.196 rows=13 loops=1)	
2	[...] Workers Planned: 3	
3	[...] Workers Launched: 3	
4	[...] -> Sort (cost=13228.04..13228.05 rows=3 width=91) (actual time=19.724..19.764 rows=3 loops=4)	
5	[...] Sort Key: mt.id	
6	[...] Sort Method: quicksort Memory: 25kB	
7	[...] Worker 0: Sort Method: quicksort Memory: 25kB	
8	[...] Worker 1: Sort Method: quicksort Memory: 25kB	
9	[...] Worker 2: Sort Method: quicksort Memory: 25kB	
10	[...] -> Nested Loop (cost=4.25..13228.02 rows=3 width=91) (actual time=7.450..19.625 rows=3 loops=4)	
11	[...] -> Hash Join (cost=3.96..13203.02 rows=3 width=53) (actual time=7.363..19.390 rows=3 loops=4)	
12	[...] Hash Cond: (mt_pl_dt.hero_id = hr.id)	
13	[...] -> Nested Loop (cost=0.42..13199.47 rows=3 width=47) (actual time=4.793..16.747 rows=3 loops=4)	
14	[...] -> Parallel Seq Scan on matches_players_details mt_pl_dt (cost=0.00..13174.13 rows=3 width=36) (actual time=4.725..16.438 rows=3 loops=4)	
15	[...] Filter: (player_id = 14944)	
16	[...] Rows Removed by Filter: 124997	
17	[...] -> Index Scan using players_pk on players pl (cost=0.42..8.44 rows=1 width=15) (actual time=0.020..0.033 rows=1 loops=13)	
18	[...] Index Cond: (id = 14944)	
19	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.407..2.419 rows=113 loops=4)	
20	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB	
21	[...] -> Seq Scan on heroes hr (cost=0.00..2.13 rows=113 width=14) (actual time=0.049..1.211 rows=113 loops=4)	
22	[...] -> Index Scan using matches_pk on matches mt (cost=0.29..8.31 rows=1 width=9) (actual time=0.026..0.028 rows=1 loops=13)	
23	[...] Index Cond: (id = mt_pl_dt.match_id)	
24	Planning Time: 0.581 ms	
25	Execution Time: 29.667 ms	

ORM generated:

```

SELECT players.id AS players_id,
CASE
    WHEN ( players.nick IS NULL ) THEN 'unknown'
    ELSE players.nick
END AS player_nick,
heroes.localized_name AS hero_localized_name,
Round(matches.duration / 60.0, 2) AS
match_duration_minutes,
COALESCE(matches_players_details.xp_hero, 0)
+ COALESCE(matches_players_details.xp_creep, 0)
+ COALESCE(matches_players_details.xp_other, 0)
+ COALESCE(matches_players_details.xp_roshan, 0) AS experiences_gained,
matches_players_details.level AS level_gained,
CASE
    WHEN ( matches_players_details.player_slot >= 0
        AND matches_players_details.player_slot <= 4
        AND matches.radiant_win = true ) THEN true
    WHEN ( matches_players_details.player_slot >= 128
        AND matches_players_details.player_slot <= 132
        AND matches.radiant_win = false ) THEN true
    ELSE false
END AS winner,
matches.id AS match_id
FROM players
JOIN matches_players_details
    ON matches_players_details.player_id = players.id
JOIN heroes
    ON matches_players_details.hero_id = heroes.id
JOIN matches
    ON matches_players_details.match_id = matches.id
WHERE players.id = 14944
ORDER BY matches.id ASC

```

	QUERY PLAN	
	text	
1	Gather Merge (cost=14228.08..14229.15 rows=9 width=91) (actual time=23.102..25.890 rows=13 loops=1)	
2	[...] Workers Planned: 3	
3	[...] Workers Launched: 3	
4	[...] -> Sort (cost=13228.04..13228.05 rows=3 width=91) (actual time=18.097..18.151 rows=3 loops=4)	
5	[...] Sort Key: matches.id	
6	[...] Sort Method: quicksort Memory: 25kB	
7	[...] Worker 0: Sort Method: quicksort Memory: 25kB	
8	[...] Worker 1: Sort Method: quicksort Memory: 25kB	
9	[...] Worker 2: Sort Method: quicksort Memory: 25kB	
10	[...] -> Nested Loop (cost=4.25..13228.02 rows=3 width=91) (actual time=9.444..18.005 rows=3 loops=4)	
11	[...] -> Hash Join (cost=3.96..13203.02 rows=3 width=53) (actual time=9.370..17.772 rows=3 loops=4)	
12	[...] Hash Cond: (matches_players_details.hero_id = heroes.id)	
13	[...] -> Nested Loop (cost=0.42..13199.47 rows=3 width=47) (actual time=7.429..15.759 rows=3 loops=4)	
14	[...] -> Parallel Seq Scan on matches_players_details (cost=0.00..13174.13 rows=3 width=36) (actual time=7.365..15.486 rows=3 loops=4)	
15	[...] Filter: (player_id = 14944)	
16	[...] Rows Removed by Filter: 124997	
17	[...] -> Index Scan using players_pk on players (cost=0.42..8.44 rows=1 width=15) (actual time=0.021..0.033 rows=1 loops=13)	
18	[...] Index Cond: (id = 14944)	
19	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.430..2.439 rows=113 loops=3)	
20	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB	
21	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.042..1.210 rows=113 loops=3)	
22	[...] -> Index Scan using matches_pk on matches (cost=0.29..8.31 rows=1 width=9) (actual time=0.026..0.028 rows=1 loops=13)	
23	[...] Index Cond: (id = matches_players_details.match_id)	
24	Planning Time: 0.558 ms	
25	Execution Time: 26.314 ms	

Summary

ORM generated absolutely the same query as original one, so it is impossible to get different performance from it.

Generally, the most expensive operations are JOIN, total cost of hash join is 13203 and the total cost of the query is 14229, while the execution time of initial query is 29.667 ms and the execution time of the ORM generated one – 26.314 ms. Such difference was caused by cache usage (the database cached some data after the initial query execution, so the ORM generated one was executed using this cached data)

3. /players/14944/game_objectives

Initial:

```
SELECT pl.id,
       CASE
         WHEN pl.nick IS NULL THEN 'unknown'
         ELSE pl.nick
       END AS player_nick,
       hr.localized_name AS hero_localized_name,
       mt.id AS match_id,
       COALESCE(gm_obj.subtype, 'NO_ACTION') AS hero_action,
       Count(*)
FROM   players AS pl
JOIN   matches_players_details AS mt_pl_dt
      ON mt_pl_dt.player_id = pl.id
JOIN   heroes AS hr
      ON mt_pl_dt.hero_id = hr.id
JOIN   matches AS mt
      ON mt_pl_dt.match_id = mt.id
LEFT JOIN game_objectives AS gm_obj
      ON mt_pl_dt.id = gm_obj.match_player_detail_id_1
WHERE  pl.id = 14944
GROUP BY pl.id,
         hr.localized_name,
         mt.id,
         gm_obj.subtype
ORDER BY match_id ASC
```

1	GroupAggregate (cost=34801.70..34802.28 rows=23 width=114) (actual time=8666.325..8666.910 rows=16 loops=1)
2	[...] Group Key: mt.id, pl.id, hr.localized_name, gm_obj.subtype
3	[...] -> Sort (cost=34801.70..34801.76 rows=23 width=53) (actual time=8666.272..8666.480 rows=18 loops=1)
4	[...] Sort Key: mt.id, hr.localized_name, gm_obj.subtype
5	[...] Sort Method: quicksort Memory: 26kB
6	[...] -> Nested Loop (cost=21591.84..34801.18 rows=23 width=53) (actual time=8652.414..8665.983 rows=18 loops=1)
7	[...] -> Index Scan using players_pk on players pl (cost=0.42..8.44 rows=1 width=15) (actual time=0.021..0.042 rows=1 loops=1)
8	[...] Index Cond: (id = 14944)
9	[...] -> Gather (cost=21591.42..34792.51 rows=23 width=42) (actual time=8652.355..8668.750 rows=18 loops=1)
10	[...] Workers Planned: 3
11	[...] Workers Launched: 3
12	[...] -> Parallel Hash Left Join (cost=20591.42..33790.21 rows=7 width=42) (actual time=8647.120..8658.855 rows=4 loops=4)
13	[...] Hash Cond: (mt_pl_dt.id = gm_obj.match_player_detail_id_1)
14	[...] -> Nested Loop (cost=3.83..13202.60 rows=3 width=22) (actual time=10.547..25.551 rows=4 loops=3)
15	[...] -> Hash Join (cost=3.54..13177.68 rows=3 width=22) (actual time=10.432..25.168 rows=4 loops=3)
16	[...] Hash Cond: (mt_pl_dt.hero_id = hr.id)
17	[...] -> Parallel Seq Scan on matches_players_details mt_pl_dt (cost=0.00..13174.13 rows=3 width=16) (actual time=6.929..21.554 rows=4 loops=3)
18	[...] Filter: (player_id = 14944)
19	[...] Rows Removed by Filter: 166662
20	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=3.398..3.409 rows=113 loops=3)
21	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
22	[...] -> Seq Scan on heroes hr (cost=0.00..2.13 rows=113 width=14) (actual time=0.060..1.707 rows=113 loops=3)
23	[...] -> Index Only Scan using matches_pk on matches mt (cost=0.29..8.31 rows=1 width=4) (actual time=0.041..0.043 rows=1 loops=13)
24	[...] Index Cond: (id = mt_pl_dt.match_id)
25	[...] Heap Fetches: 13
26	[...] -> Parallel Hash (cost=15856.15..15856.15 rows=378515 width=28) (actual time=8629.937..8629.946 rows=293349 loops=4)
27	[...] Buckets: 2097152 Batches: 1 Memory Usage: 60704kB
28	[...] -> Parallel Seq Scan on game_objectives gm_obj (cost=0.00..15856.15 rows=378515 width=28) (actual time=0.039..4339.424 rows=293349 loops=4)
29	Planning Time: 0.809 ms
30	Execution Time: 8670.962 ms

ORM generated:

```
SELECT players.id AS players_id,
       CASE
```



```

        WHEN ( players.nick IS NULL ) THEN 'unknown'
        ELSE players.nick
    END
    AS player_nick,
    heroes.localized_name
    AS hero_localized_name,
    matches.id
    AS match_id,
    COALESCE(game_objectives.subtype, 'NO_ACTION')
    AS hero_action,
    Count(*)
    AS count
FROM
    players
    JOIN matches_players_details
        ON matches_players_details.player_id = players.id
    JOIN heroes
        ON matches_players_details.hero_id = heroes.id
    JOIN matches
        ON matches_players_details.match_id = matches.id
    LEFT OUTER JOIN game_objectives
        ON matches_players_details.id =
            game_objectives.match_player_detail_id_1
WHERE
    players.id = 14944
GROUP
    BY
        players.id,
        heroes.localized_name,
        matches.id,
        game_objectives.subtype
ORDER
    BY
        matches.id ASC

```

1	GroupAggregate (cost=34801.70..34802.28 rows=23 width=114) (actual time=6131.572..6132.098 rows=16 loops=1)
2	[...] Group Key: matches.id, players.id, heroes.localized_name, game_objectives.subtype
3	[...] -> Sort (cost=34801.70..34801.76 rows=23 width=53) (actual time=6131.506..6131.681 rows=18 loops=1)
4	[...] Sort Key: matches.id, heroes.localized_name, game_objectives.subtype
5	[...] Sort Method: quicksort Memory: 26kB
6	[...] -> Nested Loop (cost=21591.84..34801.18 rows=23 width=53) (actual time=6114.820..6131.269 rows=18 loops=1)
7	[...] -> Index Scan using players_pk on players (cost=0.42..8.44 rows=1 width=15) (actual time=0.034..0.056 rows=1 loops=1)
8	[...] Index Cond: (id = 14944)
9	[...] -> Gather (cost=21591.42..34792.51 rows=23 width=42) (actual time=6114.742..6132.696 rows=18 loops=1)
10	[...] Workers Planned: 3
11	[...] Workers Launched: 3
12	[...] -> Parallel Hash Left Join (cost=20591.42..33790.21 rows=7 width=42) (actual time=6115.111..6124.173 rows=4 loops=4)
13	[...] Hash Cond: (matches_players_details.id = game_objectives.match_player_detail_id_1)
14	[...] -> Nested Loop (cost=3.83..13202.60 rows=3 width=22) (actual time=9.026..17.652 rows=3 loops=4)
15	[...] -> Hash Join (cost=3.54..13177.68 rows=3 width=22) (actual time=8.913..17.386 rows=3 loops=4)
16	[...] Hash Cond: (matches_players_details.hero_id = heroes.id)
17	[...] -> Parallel Seq Scan on matches_players_details (cost=0.00..13174.13 rows=3 width=16) (actual time=6.490..14.889 rows=3 loops=4)
18	[...] Filter: (player_id = 14944)
19	[...] Rows Removed by Filter: 124997
20	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.323..2.332 rows=113 loops=4)
21	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
22	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.056..1.154 rows=113 loops=4)
23	[...] -> Index Only Scan using matches_pk on matches (cost=0.29..8.31 rows=1 width=4) (actual time=0.040..0.043 rows=1 loops=13)
24	[...] Index Cond: (id = matches_players_details.match_id)
25	[...] Heap Fetches: 13
26	[...] -> Parallel Hash (cost=15856.15..15856.15 rows=378515 width=28) (actual time=6102.164..6102.173 rows=293349 loops=4)
27	[...] Buckets: 2097152 Batches: 1 Memory Usage: 60704kB
28	[...] -> Parallel Seq Scan on game_objectives (cost=0.00..15856.15 rows=378515 width=28) (actual time=0.041..3041.573 rows=293349 loops=4)
29	Planning Time: 1.538 ms
30	Execution Time: 6134.625 ms

Summary

ORM generated absolutely the same query as original one, so it is impossible to get different performance from it.

Generally, the most expensive operations are JOIN. The execution time of initial query is 8670 ms and the execution time of the ORM generated one – 6134 ms. Such difference was caused by cache usage (the database cached some data after the initial query execution, so the ORM generated one was executed using this cached data) Total cost is 34802 of every query.

4. /players/14944/abilities

Initial:

```
WITH table1
  AS (SELECT pl.id,
    CASE
      WHEN pl.nick IS NULL THEN 'unknown'
      ELSE pl.nick
    END AS player_nick,
    hr.localized_name AS hero_localized_name,
    mt.id AS match_id,
    ab.NAME AS ability_name,
    Max(ab_upg.level)
      OVER (
        partition BY mt.id, ab.NAME) AS upgrade_level
  FROM players AS pl
    JOIN matches_players_details AS mt_pl_dt
      ON mt_pl_dt.player_id = pl.id
    JOIN heroes AS hr
      ON mt_pl_dt.hero_id = hr.id
    JOIN matches AS mt
      ON mt_pl_dt.match_id = mt.id
    JOIN ability_upgrades AS ab_upg
      ON mt_pl_dt.id = ab_upg.match_player_detail_id
    JOIN abilities AS ab
      ON ab_upg.ability_id = ab.id
  WHERE pl.id = 14944
  GROUP BY pl.id,
    hr.localized_name,
    mt.id,
    ab.NAME,
    ab_upg.level
    ORDER BY match_id ASC)
SELECT id,
  player_nick,
  hero_localized_name,
  match_id,
  ability_name,
  upgrade_level,
  Count(*)
FROM table1
GROUP BY id,
  player_nick,
  hero_localized_name,
  match_id,
  ability_name,
  upgrade_level
ORDER BY match_id ASC
```

1	GroupAggregate (cost=99777.76..99783.13 rows=179 width=84) (actual time=37759.004..37766.384 rows=63 loops=1)
2	[...] Group Key: table1.match_id, table1.id, table1.player_nick, table1.hero_localized_name, table1.ability_name, table1.upgrade_level
3	[...] -> Sort (cost=99777.76..99778.21 rows=179 width=76) (actual time=37758.928..37762.054 rows=239 loops=1)
4	[...] Sort Key: table1.match_id, table1.id, table1.player_nick, table1.hero_localized_name, table1.ability_name, table1.upgrade_level
5	[...] Sort Method: quicksort Memory: 49kB
6	[...] -> Subquery Scan on table1 (cost=99765.69..99771.06 rows=179 width=76) (actual time=37745.065..37756.554 rows=239 loops=1)
7	[...] -> WindowAgg (cost=99765.69..99769.27 rows=179 width=80) (actual time=37745.046..37752.100 rows=239 loops=1)
8	[...] -> Sort (cost=99765.69..99766.14 rows=179 width=55) (actual time=37744.961..37747.170 rows=239 loops=1)
9	[...] Sort Key: mt.id, ab.name
10	[...] Sort Method: quicksort Memory: 49kB
11	[...] -> Group (cost=99756.31..99758.99 rows=179 width=55) (actual time=37735.358..37742.361 rows=239 loops=1)
12	[...] Group Key: mt.id, pl.id, hr.localized_name, ab.name, ab_upg.level
13	[...] -> Sort (cost=99756.31..99756.76 rows=179 width=55) (actual time=37735.324..37737.607 rows=239 loops=1)
14	[...] Sort Key: mt.id, hr.localized_name, ab.name, ab_upg.level
15	[...] Sort Method: quicksort Memory: 49kB
16	[...] -> Nested Loop (cost=14178.69..99749.61 rows=179 width=55) (actual time=13095.390..37732.522 rows=239 loops=1)
17	[...] -> Index Scan using players_pk on players pl (cost=0.42..8.44 rows=1 width=15) (actual time=0.021..0.052 rows=1 loops=1)
18	[...] Index Cond: (id = 14944)
19	[...] -> Gather (cost=14178.27..99739.38 rows=179 width=44) (actual time=13095.338..37727.796 rows=239 loops=1)
20	[...] Workers Planned: 4
21	[...] Workers Launched: 4
22	[...] -> Nested Loop (cost=13178.27..98721.48 rows=45 width=44) (actual time=19046.027..37724.405 rows=48 loops=5)
23	[...] -> Nested Loop (cost=13178.00..98708.32 rows=45 width=26) (actual time=19045.949..37722.057 rows=48 loops=5)
24	[...] -> Hash Join (cost=13177.71..98334.48 rows=45 width=26) (actual time=19045.809..37719.555 rows=48 loops=5)
25	[...] Hash Cond: (mt_pl_dt.hero_id = hr.id)
26	[...] -> Parallel Hash Join (cost=13174.17..98330.82 rows=45 width=20) (actual time=19042.653..37715.366 rows=48 loops=5)
27	[...] Hash Cond: (ab_upg.match_player_detail_id = mt_pl_dt.id)
28	[...] -> Parallel Seq Scan on ability_upgrades ab_upg (cost=0.00..79290.00 rows=2234900 width=12) (actual time=0.025..18686.349 rows=1787920 loops=5)
29	[...] -> Parallel Hash (cost=13174.13..13174.13 rows=3 width=16) (actual time=14.153..14.163 rows=3 loops=5)
30	[...] Buckets: 1024 Batches: 1 Memory Usage: 136kB
31	[...] -> Parallel Seq Scan on matches_players_details mt_pl_dt (cost=0.00..13174.13 rows=3 width=16) (actual time=5.535..13.966 rows=3 loops=5)
32	[...] Filter: (player_id = 14944)
33	[...] Rows Removed by Filter: 99997
34	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=3.002..3.012 rows=113 loops=5)
35	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
36	[...] -> Seq Scan on heroes hr (cost=0.00..2.13 rows=113 width=14) (actual time=0.043..1.508 rows=113 loops=5)
37	[...] -> Index Only Scan using matches_pk on matches mt (cost=0.29..8.31 rows=1 width=4) (actual time=0.018..0.018 rows=1 loops=239)
38	[...] Index Cond: (id = mt_pl_dt.match_id)
39	[...] Heap Fetches: 239
40	[...] -> Index Scan using abilities_pk on abilities ab (cost=0.28..0.29 rows=1 width=26) (actual time=0.016..0.016 rows=1 loops=239)
41	[...] Index Cond: (id = ab_upg.ability_id)
42	Planning Time: 1.398 ms
43	Execution Time: 37768.122 ms

ORM generated:

```
SELECT anon_1.id AS anon_1_id,
anon_1.player_nick AS anon_1_player_nick,
anon_1.hero_localized_name AS anon_1_hero_localized_name,
anon_1.match_id AS anon_1_match_id,
anon_1.ability_name AS anon_1_ability_name,
anon_1.upgrade_level AS anon_1_upgrade_level,
Count(*) AS count
FROM (SELECT players.id AS id,
CASE
WHEN ( players.nick IS NULL ) THEN 'unknown'
ELSE players.nick
END AS player_nick,
heroes.localized_name AS
hero_localized_name,
matches.id AS match_id,
abilities.NAME AS ability_name,
Max(ability_upgrades.level)
OVER (
partition BY matches.id, abilities.NAME) AS upgrade_level
FROM players
JOIN matches_players_details
ON matches_players_details.player_id = players.id
JOIN heroes
ON matches_players_details.hero_id = heroes.id
JOIN matches
ON matches_players_details.match_id = matches.id
JOIN ability_upgrades
ON matches_players_details.id =
ability_upgrades.match_player_detail_id
JOIN abilities
ON ability_upgrades.ability_id = abilities.id
WHERE players.id = 14944
GROUP BY players.id,
heroes.localized_name,
matches.id,
abilities.NAME,
ability_upgrades.level
ORDER BY matches.id ASC) AS anon_1
GROUP BY anon_1.id,
anon_1.player_nick,
anon_1.hero_localized_name,
anon_1.match_id,
anon_1.ability_name,
anon_1.upgrade_level
ORDER BY anon_1.match_id ASC
```

1	GroupAggregate (cost=99777.76..99783.13 rows=179 width=84) (actual time=37478.410..37483.983 rows=63 loops=1)
2	[...] Group Key: anon_1.match_id, anon_1.id, anon_1.player_nick, anon_1.hero_localized_name, anon_1.ability_name, anon_1.upgrade_level
3	[...] -> Sort (cost=99777.76..99778.21 rows=179 width=76) (actual time=37478.324..37480.762 rows=239 loops=1)
4	[...] Sort Key: anon_1.match_id, anon_1.id, anon_1.player_nick, anon_1.hero_localized_name, anon_1.ability_name, anon_1.upgrade_level
5	[...] Sort Method: quicksort Memory: 49kB
6	[...] -> Subquery Scan on anon_1 (cost=99765.69..99771.06 rows=179 width=76) (actual time=37464.098..37475.900 rows=239 loops=1)
7	[...] -> WindowAgg (cost=99765.69..99769.27 rows=179 width=80) (actual time=37464.078..37471.245 rows=239 loops=1)
8	[...] -> Sort (cost=99765.69..99766.14 rows=179 width=55) (actual time=37464.001..37466.221 rows=239 loops=1)
9	[...] Sort Key: matches.id, abilities.name
10	[...] Sort Method: quicksort Memory: 49kB
11	[...] -> Group (cost=99756.31..99758.99 rows=179 width=55) (actual time=37454.758..37461.460 rows=239 loops=1)
12	[...] Group Key: matches.id, players.id, heroes.localized_name, abilities.name, ability_upgrades.level
13	[...] -> Sort (cost=99756.31..99756.76 rows=179 width=55) (actual time=37454.730..37456.914 rows=239 loops=1)
14	[...] Sort Key: matches.id, heroes.localized_name, abilities.name, ability_upgrades.level
15	[...] Sort Method: quicksort Memory: 49kB
16	[...] -> Nested Loop (cost=14178.69..99749.61 rows=179 width=55) (actual time=13650.112..37452.019 rows=239 loops=1)
17	[...] -> Index Scan using players_pk on players (cost=0.42..8.44 rows=1 width=15) (actual time=0.022..0.053 rows=1 loops=1)
18	[...] Index Cond: (id = 14944)
19	[...] -> Gather (cost=14178.27..99739.38 rows=179 width=44) (actual time=13650.058..37447.340 rows=239 loops=1)
20	[...] Workers Planned: 4
21	[...] Workers Launched: 4
22	[...] -> Nested Loop (cost=13178.27..98721.48 rows=45 width=44) (actual time=15970.125..37444.745 rows=48 loops=5)
23	[...] -> Nested Loop (cost=13178.00..98708.32 rows=45 width=26) (actual time=15970.044..37442.420 rows=48 loops=5)
24	[...] -> Hash Join (cost=13177.71..98334.48 rows=45 width=26) (actual time=15969.923..37439.917 rows=48 loops=5)
25	[...] Hash Cond: (matches_players_details.hero_id = heroes.id)
26	[...] -> Parallel Hash Join (cost=13174.17..98330.82 rows=45 width=20) (actual time=15967.249..37436.213 rows=48 loops=5)
27	[...] Hash Cond: (ability_upgrades.match_player_detail_id = matches_players_details.id)
28	[...] -> Parallel Seq Scan on ability_upgrades (cost=0.00..79290.00 rows=2234900 width=12) (actual time=0.027..18544.646 rows=1787920 loops=5)
29	[...] -> Parallel Hash (cost=13174.13..13174.13 rows=3 width=16) (actual time=12.835..12.844 rows=3 loops=5)
30	[...] Buckets: 1024 Batches: 1 Memory Usage: 168kB
31	[...] -> Parallel Seq Scan on matches_players_details (cost=0.00..13174.13 rows=3 width=16) (actual time=5.649..12.669 rows=3 loops=5)
32	[...] Filter: (player_id = 14944)
33	[...] Rows Removed by Filter: 99997
34	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.529..2.538 rows=113 loops=5)
35	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
36	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.041..1.285 rows=113 loops=5)
37	[...] -> Index Only Scan using matches_pk on matches (cost=0.29..8.31 rows=1 width=4) (actual time=0.018..0.018 rows=1 loops=239)
38	[...] Index Cond: (id = matches_players_details.match_id)
39	[...] Heap Fetches: 239
40	[...] -> Index Scan using abilities_pk on abilities (cost=0.28..0.29 rows=1 width=26) (actual time=0.015..0.015 rows=1 loops=239)
41	[...] Index Cond: (id = ability_upgrades.ability_id)
42	Planning Time: 1.397 ms
43	Execution Time: 37485.279 ms

Summary

ORM generated query is similar to the original one, so it is impossible to get different performance from it. The only difference is that in original query was used WITH to make subtable, but in the generated query was used nested SELECT. However, query plan has the same structure and almost the same values in both cases.

Generally, the most expensive operations are JOIN. The execution time of initial query is 37768 ms and the execution time of the ORM generated one – 37485 ms. The difference in execution time can be considered as statistical deviation. Total cost is 99783 of every query.

5. /matches/21421/top_purchases/

Initial:

```
SELECT *
FROM (
    SELECT *,
           Row_number() OVER (partition BY localized_name ORDER BY hero_id ASC, count
DESC, item_name ASC) AS row_num
    FROM (
        SELECT m.id AS match_id,
               m_p_d.hero_id,
               her.localized_name,
               p_log.item_id,
               itms.NAME AS item_name,
               Count(*)
        FROM purchase_logs AS p_log
        JOIN matches_players_details AS m_p_d
        ON p_log.match_player_detail_id = m_p_d.id
        JOIN heroes AS her
        ON m_p_d.hero_id = her.id
        JOIN matches AS m
        ON m_p_d.match_id = m.id
        JOIN items AS itms
        ON p_log.item_id = itms.id
        WHERE ((
            AND m_p_d.player_slot >= 0
            AND m_p_d.player_slot <= 4
            AND m.radiant_win IS true)
            OR (
            AND m_p_d.player_slot >= 128
            AND m_p_d.player_slot <= 132
            AND m.radiant_win IS false))
        AND m.id = 21421
        GROUP BY m.id,
               m_p_d.hero_id,
               her.localized_name,
               p_log.item_id,
               itms.NAME) table1
    ORDER BY hero_id ASC,
           count DESC,
           item_name ASC) table2
WHERE row_num <= 5
```


1	Subquery Scan on table2 (cost=156879.44..156882.41 rows=66 width=52) (actual time=80571.414..80575.520 rows=25 loops=1)
2	[...] Filter: (table2.row_num <= 5)
3	[...] Rows Removed by Filter: 88
4	[...] -> Sort (cost=156879.44..156879.93 rows=198 width=52) (actual time=80571.387..80573.231 rows=113 loops=1)
5	[...] Sort Key: table1.hero_id, table1.count DESC, table1.item_name
6	[...] Sort Method: quicksort Memory: 37kB
7	[...] -> WindowAgg (cost=156866.94..156871.89 rows=198 width=52) (actual time=80563.946..80569.384 rows=113 loops=1)
8	[...] -> Sort (cost=156866.94..156867.43 rows=198 width=44) (actual time=80563.872..80565.721 rows=113 loops=1)
9	[...] Sort Key: table1.localized_name, table1.hero_id, table1.count DESC, table1.item_name
10	[...] Sort Method: quicksort Memory: 34kB
11	[...] -> Subquery Scan on table1 (cost=156851.96..156859.38 rows=198 width=44) (actual time=80551.134..80561.730 rows=113 loops=1)
12	[...] -> GroupAggregate (cost=156851.96..156857.40 rows=198 width=44) (actual time=80551.115..80558.470 rows=113 loops=1)
13	[...] Group Key: m.id, m_p.d.hero_id, her.localized_name, p_log.item_id, itms.name
14	[...] -> Sort (cost=156851.96..156852.45 rows=198 width=36) (actual time=80551.049..80553.674 rows=188 loops=1)
15	[...] Sort Key: m_p.d.hero_id, her.localized_name, p_log.item_id, itms.name
16	[...] Sort Method: quicksort Memory: 39kB
17	[...] -> Nested Loop (cost=1036.40..156844.40 rows=198 width=36) (actual time=35543.823..80548.969 rows=188 loops=1)
18	[...] Join Filter: (((m_p.d.player_slot >= 0) AND (m_p.d.player_slot <= 4) AND (m.radiant_win IS TRUE)) OR (((m_p.d.player_slot >= 128) AND (m_p.d.player_slot <= 132) AND (m.radiant_win IS FALSE)))
19	[...] Rows Removed by Join Filter: 165
20	[...] -> Gather (cost=1036.11..156831.06 rows=255 width=40) (actual time=35543.598..80533.533 rows=353 loops=1)
21	[...] Workers Planned: 4
22	[...] Workers Launched: 4
23	[...] -> Hash Join (cost=36.11..155805.56 rows=64 width=40) (actual time=53539.918..80538.094 rows=71 loops=5)
24	[...] Hash Cond: (p_log.item_id = itms.id)
25	[...] -> Hash Join (cost=28.08..155797.36 rows=64 width=26) (actual time=53534.386..80531.145 rows=71 loops=5)
26	[...] Hash Cond: (m_p.d.hero_id = her.id)
27	[...] -> Hash Join (cost=24.54..155793.64 rows=64 width=16) (actual time=53532.013..80527.382 rows=71 loops=5)
28	[...] Hash Cond: (p_log.match_player_detail_id = m_p.d.id)
29	[...] -> Parallel Seq Scan on purchase_logs p_log (cost=0.00..143829.36 rows=4548436 width=8) (actual time=0.019..40225.434 rows=3638749 loops=5)
30	[...] -> Hash (cost=24.45..24.45 rows=7 width=16) (actual time=0.300..0.312 rows=10 loops=5)
31	[...] Buckets: 1024 Batches: 1 Memory Usage: 9kB
32	[...] -> Index Scan using idx_match_id_player_id on matches_players_details m_p.d (cost=0.42..24.45 rows=7 width=16) (actual time=0.035..0.185 rows=10 loops=5)
33	[...] Index Cond: (match_id = 21421)
34	[...] Filter: (((player_slot >= 0) AND (player_slot <= 4)) OR ((player_slot >= 128) AND (player_slot <= 132)))
35	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.308..2.320 rows=113 loops=5)
36	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
37	[...] -> Seq Scan on heroes her (cost=0.00..2.13 rows=113 width=14) (actual time=0.029..1.168 rows=113 loops=5)
38	[...] -> Hash (cost=4.68..4.68 rows=268 width=18) (actual time=5.407..5.420 rows=268 loops=5)
39	[...] Buckets: 1024 Batches: 1 Memory Usage: 22kB
40	[...] -> Seq Scan on items itms (cost=0.00..4.68 rows=268 width=18) (actual time=0.040..2.714 rows=268 loops=5)
41	[...] -> Materialize (cost=0.29..8.31 rows=1 width=5) (actual time=0.010..0.014 rows=1 loops=353)
42	[...] -> Index Scan using matches_pk on matches m (cost=0.29..8.31 rows=1 width=5) (actual time=0.041..0.063 rows=1 loops=1)
43	[...] Index Cond: (id = 21421)
44	[...] Filter: ((radiant_win IS TRUE) OR (radiant_win IS FALSE))
45	Planning Time: 1.078 ms
46	Execution Time: 80576.700 ms

ORM generated:

```

SELECT anon_1.match_id           AS anon_1_match_id,
       anon_1.hero_id           AS anon_1_hero_id,
       anon_1.localized_name     AS anon_1_localized_name,
       anon_1.item_id           AS anon_1_item_id,
       anon_1.item_name         AS anon_1_item_name,
       anon_1.count             AS anon_1_count,
       anon_1.row_num           AS anon_1_row_num
FROM   (SELECT anon_2.match_id
        AS
            match_id,
        anon_2.hero_id
        AS hero_id,
        anon_2.localized_name
        AS localized_name,
        anon_2.item_id
        AS item_id,
        anon_2.item_name
        AS item_name,
        anon_2.count
        AS count,
        Row_number ()
        OVER (

```

```

        partition BY anon_2.localized_name
        ORDER BY anon_2.hero_id ASC, anon_2.count DESC,
        anon_2.item_name ASC
    ) AS
row_num
FROM (SELECT matches.id AS match_id,
        matches_players_details.hero_id AS hero_id,
        heroes.localized_name AS localized_name,
        purchase_logs.item_id AS item_id,
        items.NAME AS item_name,
        Count(*) AS count
    FROM purchase_logs
    JOIN matches_players_details
        ON purchase_logs.match_player_detail_id =
            matches_players_details.id
    JOIN heroes
        ON matches_players_details.hero_id = heroes.id
    JOIN matches
        ON matches_players_details.match_id = matches.id
    JOIN items
        ON purchase_logs.item_id = items.id
    WHERE ( matches_players_details.player_slot >= 0
        AND matches_players_details.player_slot <= 4
        AND matches.radiant_win = true
        OR matches_players_details.player_slot >= 128
        AND matches_players_details.player_slot <= 132
        AND matches.radiant_win = false )
        AND matches.id = 21421
    GROUP BY matches.id,
        matches_players_details.hero_id,
        heroes.localized_name,
        purchase_logs.item_id,
        items.NAME) AS anon_2
ORDER BY anon_2.hero_id ASC,
        anon_2.count DESC,
        anon_2.item_name ASC) AS anon_1
WHERE anon_1.row_num <= 5

```

1	Subquery Scan on anon_1 (cost=156879.44..156882.41 rows=66 width=52) (actual time=79008.151..79010.480 rows=25 loops=1)
2	[...] Filter: (anon_1.row_num <= 5)
3	[...] Rows Removed by Filter: 88
4	[...] Sort (cost=156879.44..156879.93 rows=198 width=52) (actual time=79008.123..79009.191 rows=113 loops=1)
5	[...] Sort Key: anon_2.hero_id, anon_2.count DESC, anon_2.item_name
6	[...] Sort Method: quicksort Memory: 37kB
7	[...] WindowAgg (cost=156866.94..156871.89 rows=198 width=52) (actual time=79003.412..79006.856 rows=113 loops=1)
8	[...] Sort (cost=156866.94..156867.43 rows=198 width=44) (actual time=79003.348..79004.434 rows=113 loops=1)
9	[...] Sort Key: anon_2.localized_name, anon_2.hero_id, anon_2.count DESC, anon_2.item_name
10	[...] Sort Method: quicksort Memory: 34kB
11	[...] Subquery Scan on anon_2 (cost=156851.96..156859.38 rows=198 width=44) (actual time=78995.140..79002.082 rows=113 loops=1)
12	[...] GroupAggregate (cost=156851.96..156857.40 rows=198 width=44) (actual time=78995.115..78999.920 rows=113 loops=1)
13	[...] Group Key: matches.id, matches.players_details.hero_id, heroes.localized_name, purchase_logs.item_id, items.name
14	[...] Sort (cost=156851.96..156852.45 rows=198 width=36) (actual time=78995.062..78996.850 rows=188 loops=1)
15	[...] Sort Key: matches.players_details.hero_id, heroes.localized_name, purchase_logs.item_id, items.name
16	[...] Sort Method: quicksort Memory: 39kB
17	[...] Nested Loop (cost=1036.40..156844.40 rows=198 width=36) (actual time=78974.088..78993.032 rows=188 loops=1)
18	[...] Join Filter: (((matches.players_details.player_slot >= 0) AND (matches.players_details.player_slot <= 4) AND matches.radiant_win) OR ((matches.players_details.player_slot >= 128) AND (matches.players_details.player_slot <= 132) AND (NOT matches.radiant_win)))
19	[...] Rows Removed by Join Filter: 165
20	[...] Gather (cost=1036.11..156831.06 rows=255 width=40) (actual time=78973.967..78977.564 rows=353 loops=1)
21	[...] Workers Planned: 4
22	[...] Workers Launched: 4
23	[...] Hash Join (cost=36.11..155805.56 rows=64 width=40) (actual time=51782.093..78969.177 rows=71 loops=5)
24	[...] Hash Cond: (purchase_logs.item_id = items.id)
25	[...] Hash Join (cost=28.08..155797.36 rows=64 width=26) (actual time=51775.774..78961.422 rows=71 loops=5)
26	[...] Hash Cond: (matches.players_details.hero_id = heroes.id)
27	[...] Hash Join (cost=24.54..155793.64 rows=64 width=16) (actual time=51773.056..78957.285 rows=71 loops=5)
28	[...] Hash Cond: (purchase_logs.match_player_detail_id = matches.players_details.id)
29	[...] Parallel Seq Scan on purchase_logs (cost=0.00..143829.36 rows=4548436 width=8) (actual time=0.021..39400.331 rows=3638749 loops=5)
30	[...] Hash (cost=24.45..24.45 rows=7 width=16) (actual time=0.313..0.328 rows=10 loops=5)
31	[...] Buckets: 1024 Batches: 1 Memory Usage: 9kB
32	[...] Index Scan using idx_match_id_player_id on matches.players_details (cost=0.42..24.45 rows=7 width=16) (actual time=0.037..0.171 rows=10 loops=5)
33	[...] Index Cond: (match_id = 21421)
34	[...] Filter: (((player_slot >= 0) AND (player_slot <= 4)) OR ((player_slot >= 128) AND (player_slot <= 132)))
35	[...] Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.642..2.652 rows=113 loops=5)
36	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
37	[...] Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.031..1.319 rows=113 loops=5)
38	[...] Hash (cost=4.68..4.68 rows=268 width=18) (actual time=6.186..6.195 rows=268 loops=5)
39	[...] Buckets: 1024 Batches: 1 Memory Usage: 22kB
40	[...] Seq Scan on items (cost=0.00..4.68 rows=268 width=18) (actual time=0.036..3.111 rows=268 loops=5)
41	[...] Materialize (cost=0.29..8.31 rows=1 width=5) (actual time=0.010..0.014 rows=1 loops=353)
42	[...] Index Scan using matches_pk on matches (cost=0.29..8.31 rows=1 width=5) (actual time=0.045..0.071 rows=1 loops=1)
43	[...] Index Cond: (id = 21421)
44	[...] Filter: (radiant_win OR (NOT radiant_win))
45	Planning Time: 1.887 ms
46	Execution Time: 79011.339 ms

Summary

ORM generated query is identical to the original one, so it is impossible to get different performance from it.

Generally, the most expensive operations are JOIN. The execution time of initial query is 80576 ms and the execution time of the ORM generated one – 79011 ms. The difference in execution time can be considered as statistical deviation. Total cost is 156882 of every query.

6. /abilities/5004/usage/

Initial:

```
SELECT DISTINCT
ON (
    hero_id, winner) ability_id,
    NAME,
    hero_id,
    hero_name,
    winner,
    Concat(Substr(percent_time, 1, Length(percent_time) - 1), '0-
', Substr(percent_time, 1, Length(percent_time) - 1), '9') AS rozsah,
    Count(*)
FROM
    (
        SELECT ab_up.ability_id,
            ab.NAME,
            her.id AS hero_id,
            her.localized_name AS hero_name,
            CASE
                WHEN m_p_d.player_slot >= 0
                AND m_p_d.player_slot <= 4
                AND m.radiant_win IS true THEN true
                WHEN m_p_d.player_slot >= 128
                AND m_p_d.player_slot <= 132
                AND m.radiant_win IS false THEN true
                ELSE false
            END AS winner,
            cast((ab_up.time * 100 / m.duration) AS varchar) AS percent_time
        FROM ability_upgrades AS ab_up
        JOIN abilities AS ab
        ON ab.id = ab_up.ability_id
        JOIN matches_players_details AS m_p_d
        ON ab_up.match_player_detail_id = m_p_d.id
        JOIN heroes AS her
        ON her.id = m_p_d.hero_id
        JOIN matches AS m
        ON m.id = m_p_d.match_id
        WHERE ab_up.ability_id = 5004) AS table1
GROUP BY
    ability_id,
    NAME,
    hero_id,
    hero_name,
    winner,
    rozsah
ORDER BY
    hero_id ASC,
    winner DESC,
    count DESC
```

1	Unique (cost=115593.91..115879.98 rows=38142 width=81) (actual time=6171.420..6172.101 rows=3 loops=1)
2	[...] -> Sort (cost=115593.91..115689.27 rows=38142 width=81) (actual time=6171.399..6171.725 rows=34 loops=1)
3	[...] Sort Key: her.id, (CASE WHEN ((m_p_d.player_slot >= 0) AND (m_p_d.player_slot <= 4) AND (m.radiant_win IS TRUE)) THEN true WHEN ((m_p_d.player_slot >= 128) AND (m_
4	[...] Sort Method: quicksort Memory: 29kB
5	[...] -> HashAggregate (cost=109735.48..112691.48 rows=38142 width=81) (actual time=6170.456..6171.022 rows=34 loops=1)
6	[...] Group Key: ab_up.ability_id, ab.name, her.id, CASE WHEN ((m_p_d.player_slot >= 0) AND (m_p_d.player_slot <= 4) AND (m.radiant_win IS TRUE)) THEN true WHEN ((m_p_d.p
7	[...] -> Nested Loop (cost=17431.85..109163.35 rows=38142 width=73) (actual time=4095.025..5680.302 rows=37068 loops=1)
8	[...] -> Index Scan using abilities_pk on abilities ab (cost=0.28..8.29 rows=1 width=26) (actual time=0.020..0.047 rows=1 loops=1)
9	[...] Index Cond: (id = 5004)
10	[...] -> Gather (cost=17431.57..106199.05 rows=38142 width=31) (actual time=4094.936..4701.195 rows=37068 loops=1)
11	[...] Workers Planned: 4
12	[...] Workers Launched: 4
13	[...] -> Hash Join (cost=16431.57..101384.85 rows=9536 width=31) (actual time=4098.314..4906.661 rows=7414 loops=5)
14	[...] Hash Cond: (m_p_d.match_id = m.id)
15	[...] -> Hash Join (cost=14790.57..99718.82 rows=9536 width=30) (actual time=2689.824..3333.870 rows=7414 loops=5)
16	[...] Hash Cond: (m_p_d.hero_id = her.id)
17	[...] -> Parallel Hash Join (cost=14787.03..99689.31 rows=9536 width=20) (actual time=2686.901..3162.270 rows=7414 loops=5)
18	[...] Hash Cond: (ab_up.match_player_detail_id = m_p_d.id)
19	[...] -> Parallel Seq Scan on ability_upgrades ab_up (cost=0.00..84877.25 rows=9536 width=12) (actual time=0.225..273.528 rows=7414 loops=5)
20	[...] Filter: (ability_id = 5004)
21	[...] Rows Removed by Filter: 1780506
22	[...] -> Parallel Hash (cost=12770.90..12770.90 rows=161290 width=16) (actual time=2676.390..2676.399 rows=100000 loops=5)
23	[...] Buckets: 524288 Batches: 1 Memory Usage: 27648kB
24	[...] -> Parallel Seq Scan on matches_players_details m_p_d (cost=0.00..12770.90 rows=161290 width=16) (actual time=0.038..1259.065 rows=100000 loops=5)
25	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.830..2.840 rows=113 loops=5)
26	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
27	[...] -> Seq Scan on heroes her (cost=0.00..2.13 rows=113 width=14) (actual time=0.055..1.413 rows=113 loops=5)
28	[...] -> Hash (cost=1016.00..1016.00 rows=50000 width=9) (actual time=1407.865..1407.874 rows=50000 loops=5)
29	[...] Buckets: 65536 Batches: 1 Memory Usage: 2661kB
30	[...] -> Seq Scan on matches m (cost=0.00..1016.00 rows=50000 width=9) (actual time=0.039..705.985 rows=50000 loops=5)
31	Planning Time: 0.955 ms
32	Execution Time: 6173.752 ms

ORM generated:

```

SELECT DISTINCT ON ( anon_1.hero_id, anon_1.winner)
anon_1.ability_id
anon_1_ability_id,
anon_1.NAME
AS anon_1_name,
anon_1.hero_id
AS anon_1_hero_id,
anon_1.hero_name
AS anon_1_hero_name,
anon_1.winner
AS anon_1_winner,
Concat(
Substr(anon_1.percent_time, 1,
Length(anon_1.percent_time) - 1), '0-',
Substr(anon_1.percent_time, 1,
Length(anon_1.percent_time) - 1), '9') AS rozisah,
Count(*)
AS
count
FROM (SELECT ability_upgrades.ability_id
AS
ability_id,
abilities.NAME
AS
NAME,
heroes.id
AS
hero_id,
heroes.localized_name
AS

```

```

        hero_name,
CASE
    WHEN ( matches_players_details.player_slot >= 0
          AND matches_players_details.player_slot <= 4
          AND matches.radiant_win = true ) THEN true
    WHEN ( matches_players_details.player_slot >= 128
          AND matches_players_details.player_slot <= 132
          AND matches.radiant_win = false ) THEN true
    ELSE false
END
AS
    winner,
Cast(( ability_upgrades.time * 100 ) / matches.duration AS
      VARCHAR) AS
    percent_time
FROM
    ability_upgrades
JOIN
    abilities
    ON abilities.id = ability_upgrades.ability_id
JOIN
    matches_players_details
    ON ability_upgrades.match_player_detail_id =
       matches_players_details.id
JOIN
    heroes
    ON heroes.id = matches_players_details.hero_id
JOIN
    matches
    ON matches.id = matches_players_details.match_id
WHERE
    ability_upgrades.ability_id = 5004) AS anon_1
GROUP BY
    anon_1.ability_id,
    anon_1.NAME,
    anon_1.hero_id,
    anon_1.hero_name,
    anon_1.winner,
    rozsah
ORDER BY
    anon_1.hero_id ASC,
    anon_1.winner DESC,
    Count(*) DESC

```

1	Unique (cost=115593.91..115879.98 rows=38142 width=81) (actual time=5683.688..5684.369 rows=3 loops=1)
2	[...] -> Sort (cost=115593.91..115689.27 rows=38142 width=81) (actual time=5683.663..5683.992 rows=34 loops=1)
3	[...] Sort Key: heroes.id, (CASE WHEN ((matches_players_details.player_slot >= 0) AND (matches_players_details.player_slot <= 4) AND matches.radiant_win) THE
4	[...] Sort Method: quicksort Memory: 29kB
5	[...] -> HashAggregate (cost=109735.48..112691.48 rows=38142 width=81) (actual time=5682.666..5683.264 rows=34 loops=1)
6	[...] Group Key: ability_upgrades.ability_id, abilities.name, heroes.id, CASE WHEN ((matches_players_details.player_slot >= 0) AND (matches_players_details.playe
7	[...] -> Nested Loop (cost=17431.85..109163.35 rows=38142 width=73) (actual time=3773.402..5208.659 rows=37068 loops=1)
8	[...] -> Index Scan using abilities_pk on abilities (cost=0.28..8.29 rows=1 width=26) (actual time=0.026..0.048 rows=1 loops=1)
9	[...] Index Cond: (id = 5004)
10	[...] -> Gather (cost=17431.57..106199.05 rows=38142 width=31) (actual time=3773.311..4298.772 rows=37068 loops=1)
11	[...] Workers Planned: 4
12	[...] Workers Launched: 4
13	[...] -> Hash Join (cost=16431.57..101384.85 rows=9536 width=31) (actual time=3789.464..4701.823 rows=7414 loops=5)
14	[...] Hash Cond: (matches_players_details.match_id = matches.id)
15	[...] -> Hash Join (cost=14790.57..99718.82 rows=9536 width=30) (actual time=2559.020..3291.747 rows=7414 loops=5)
16	[...] Hash Cond: (matches_players_details.hero_id = heroes.id)
17	[...] -> Parallel Hash Join (cost=14787.03..99689.31 rows=9536 width=20) (actual time=2556.555..3084.755 rows=7414 loops=5)
18	[...] Hash Cond: (ability_upgrades.match_player_detail_id = matches_players_details.id)
19	[...] -> Parallel Seq Scan on ability_upgrades (cost=0.00..84877.25 rows=9536 width=12) (actual time=0.155..343.324 rows=7414 loops=5)
20	[...] Filter: (ability_id = 5004)
21	[...] Rows Removed by Filter: 1780506
22	[...] -> Parallel Hash (cost=12770.90..12770.90 rows=161290 width=16) (actual time=2527.334..2527.344 rows=100000 loops=5)
23	[...] Buckets: 524288 Batches: 1 Memory Usage: 27616kB
24	[...] -> Parallel Seq Scan on matches_players_details (cost=0.00..12770.90 rows=161290 width=16) (actual time=0.026..1232.559 rows=100000 loops=5)
25	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.391..2.401 rows=113 loops=5)
26	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
27	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.037..1.200 rows=113 loops=5)
28	[...] -> Hash (cost=1016.00..1016.00 rows=50000 width=9) (actual time=1230.025..1230.037 rows=50000 loops=5)
29	[...] Buckets: 65536 Batches: 1 Memory Usage: 2661kB
30	[...] -> Seq Scan on matches (cost=0.00..1016.00 rows=50000 width=9) (actual time=0.030..613.422 rows=50000 loops=5)
31	Planning Time: 1.478 ms
32	Execution Time: 5686.164 ms

Summary

ORM generated query is absolutely identical to the original one, so it is impossible to get different performance from it.

Generally, the most expensive operations are JOIN. The execution time of initial query is 6173 ms and the execution time of the ORM generated one – 5686 ms. The difference in execution time can be considered as statistical deviation. Total cost is 115879 of every query.

7. /statistics/tower_kills/

Initial:

```
SELECT hero_id,
       heroes.localized_name,
       max
FROM   (SELECT hero_id,
               Max(count)
       FROM   (SELECT match_id,
                       hero_id,
                       Count(*)
               FROM   (SELECT Row_number()
                       OVER (
                           ORDER BY m_p_d.match_id ASC, g_ob.time ASC)
                       - Row_number()
                       OVER (
                           partition BY m_p_d.match_id,
                                       m_p_d.hero_id
                           ORDER BY m_p_d.match_id ASC,
                                       g_ob.time
                                       ASC) AS
               row_num,
               m_p_d.match_id,
               m_p_d.hero_id,
               g_ob.time
       FROM   game_objectives AS g_ob
       JOIN   matches_players_details AS m_p_d
               ON g_ob.match_player_detail_id_1 = m_p_d.id
       WHERE  g_ob.subtype = 'CHAT_MESSAGE_TOWER_KILL'
               AND match_player_detail_id_1 IS NOT NULL
       ORDER BY m_p_d.match_id ASC,
               g_ob.time ASC) table1
       GROUP BY row_num,
               match_id,
               hero_id) table2
       GROUP BY hero_id) table3
JOIN   heroes
       ON hero_id = heroes.id
ORDER BY max DESC,
       localized_name ASC
```


1	Sort (cost=154372.02..154372.30 rows=113 width=22) (actual time=74401.741..74402.752 rows=110 loops=1)
2	[...] Sort Key: (max((count(*)))) DESC, heroes.localized_name
3	[...] Sort Method: quicksort Memory: 33kB
4	[...] -> Hash Join (cost=154363.63..154368.16 rows=113 width=22) (actual time=74397.471..74400.601 rows=110 loops=1)
5	[...] Hash Cond: (m_p_d.hero_id = heroes.id)
6	[...] -> HashAggregate (cost=154360.08..154362.08 rows=200 width=12) (actual time=74395.130..74396.187 rows=110 loops=1)
7	[...] Group Key: m_p_d.hero_id
8	[...] -> HashAggregate (cost=153337.56..153746.57 rows=40901 width=24) (actual time=66964.124..70706.248 rows=375232 loops=1)
9	[...] Group Key: (row_number() OVER (?) - (row_number() OVER (?))), m_p_d.match_id, m_p_d.hero_id
10	[...] -> WindowAgg (cost=136070.32..145211.80 rows=406288 width=20) (actual time=47862.636..61933.743 rows=475710 loops=1)
11	[...] -> Sort (cost=136070.32..137086.04 rows=406288 width=20) (actual time=47862.578..52345.890 rows=475710 loops=1)
12	[...] Sort Key: m_p_d.match_id, g_ob."time"
13	[...] Sort Method: quicksort Memory: 49453kB
14	[...] -> WindowAgg (cost=88063.04..98220.24 rows=406288 width=20) (actual time=28618.384..43004.493 rows=475710 loops=1)
15	[...] -> Sort (cost=88063.04..89078.76 rows=406288 width=12) (actual time=28618.331..33170.081 rows=475710 loops=1)
16	[...] Sort Key: m_p_d.match_id, m_p_d.hero_id, g_ob."time"
17	[...] Sort Method: quicksort Memory: 34587kB
18	[...] -> Hash Join (cost=22408.00..50212.96 rows=406288 width=12) (actual time=9819.090..23837.236 rows=475710 loops=1)
19	[...] Hash Cond: (g_ob.match_player_detail_id_1 = m_p_d.id)
20	[...] -> Seq Scan on game_objectives g_ob (cost=0.00..26738.45 rows=406288 width=8) (actual time=0.021..4863.542 rows=475710 loops=1)
21	[...] Filter: ((match_player_detail_id_1 IS NOT NULL) AND (subtype = 'CHAT_MESSAGE_TOWER_KILL::text'))
22	[...] Rows Removed by Filter: 697686
23	[...] -> Hash (cost=16158.00..16158.00 rows=500000 width=12) (actual time=9818.479..9818.488 rows=500000 loops=1)
24	[...] Buckets: 524288 Batches: 1 Memory Usage: 25581kB
25	[...] -> Seq Scan on matches_players_details m_p_d (cost=0.00..16158.00 rows=500000 width=12) (actual time=0.018..4862.904 rows=500000 loops=1)
26	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.284..2.293 rows=113 loops=1)
27	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
28	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.023..1.158 rows=113 loops=1)
29	Planning Time: 0.519 ms
30	Execution Time: 74414.133 ms

ORM generated:

```

SELECT anon_1.hero_id AS anon_1_hero_id,
       heroes.localized_name AS heroes_localized_name,
       anon_1.max_1 AS anon_1_max_1
FROM (
    SELECT anon_2.hero_id AS hero_id,
           Max(anon_2.count) AS max_1
    FROM (
        SELECT anon_3.match_id AS match_id,
               anon_3.hero_id AS hero_id,
               Count(*) AS count
        FROM (
            SELECT Row_number() OVER (ORDER BY matches_p
layers_details.match_id ASC, game_objectives.time ASC) - Row_number() OVER (partition BY matches_pla
yers_details.match_id, matches_players_details.hero_id ORDER BY matches_players_details.match_id ASC
, game_objectives.time ASC) AS row_num,
                                           matches_players_details.match_id
                                           AS ma
tch_id,
                                           matches_players_details.hero_id
                                           AS her
o_id,
                                           game_objectives.time
                                           AS time
        FROM game_objectives
        JOIN matches_players_details
        ON game_objectives.match_player_detail_i

```

```

d_1 = matches_players_details.id
GE_TOWER_KILL'
d_1 IS NOT NULL

WHERE game_objectives.subtype = 'CHAT_MESSA
AND game_objectives.match_player_detail_i
ORDER BY matches_players_details.match_id ASC,
game_objectives.time ASC) AS anon_3
GROUP BY anon_3.row_num,
anon_3.match_id,
anon_3.hero_id) AS anon_2
GROUP BY anon_2.hero_id) AS anon_1
JOIN heroes
ON anon_1.hero_id = heroes.id
ORDER BY anon_1.max_1 DESC,
heroes.localized_name ASC

```

1	Sort (cost=154372.02..154372.30 rows=113 width=22) (actual time=74236.698..74237.704 rows=110 loops=1)
2	[...] Sort Key: (max((count(*)))) DESC, heroes.localized_name
3	[...] Sort Method: quicksort Memory: 33kB
4	[...] -> Hash Join (cost=154363.63..154368.16 rows=113 width=22) (actual time=74232.408..74235.556 rows=110 loops=1)
5	[...] Hash Cond: (matches_players_details.hero_id = heroes.id)
6	[...] -> HashAggregate (cost=154360.08..154362.08 rows=200 width=12) (actual time=74230.088..74231.133 rows=110 loops=1)
7	[...] Group Key: matches_players_details.hero_id
8	[...] -> HashAggregate (cost=153337.56..153746.57 rows=40901 width=24) (actual time=66822.224..70534.322 rows=375232 loops=1)
9	[...] Group Key: (row_number() OVER (?) - (row_number() OVER (?))), matches_players_details.match_id, matches_players_details.hero_id
10	[...] -> WindowAgg (cost=136070.32..145211.80 rows=406288 width=20) (actual time=47719.200..61799.598 rows=475710 loops=1)
11	[...] -> Sort (cost=136070.32..137086.04 rows=406288 width=20) (actual time=47719.147..52240.456 rows=475710 loops=1)
12	[...] Sort Key: matches_players_details.match_id, game_objectives.time
13	[...] Sort Method: quicksort Memory: 49453kB
14	[...] -> WindowAgg (cost=88063.04..98220.24 rows=406288 width=20) (actual time=28690.451..42879.486 rows=475710 loops=1)
15	[...] -> Sort (cost=88063.04..89078.76 rows=406288 width=12) (actual time=28690.399..33179.427 rows=475710 loops=1)
16	[...] Sort Key: matches_players_details.match_id, matches_players_details.hero_id, game_objectives.time
17	[...] Sort Method: quicksort Memory: 34587kB
18	[...] -> Hash Join (cost=22408.00..50212.96 rows=406288 width=12) (actual time=9905.236..23919.604 rows=475710 loops=1)
19	[...] Hash Cond: (game_objectives.match_player_detail_id_1 = matches_players_details.id)
20	[...] -> Seq Scan on game_objectives (cost=0.00..26738.45 rows=406288 width=8) (actual time=0.020..4803.601 rows=475710 loops=1)
21	[...] Filter: ((match_player_detail_id_1 IS NOT NULL) AND (subtype = 'CHAT_MESSAGE_TOWER_KILL::text))
22	[...] Rows Removed by Filter: 697686
23	[...] -> Hash (cost=16158.00..16158.00 rows=500000 width=12) (actual time=9904.645..9904.654 rows=500000 loops=1)
24	[...] Buckets: 524288 Batches: 1 Memory Usage: 25581kB
25	[...] -> Seq Scan on matches_players_details (cost=0.00..16158.00 rows=500000 width=12) (actual time=0.018..4890.741 rows=500000 loops=1)
26	[...] -> Hash (cost=2.13..2.13 rows=113 width=14) (actual time=2.283..2.292 rows=113 loops=1)
27	[...] Buckets: 1024 Batches: 1 Memory Usage: 14kB
28	[...] -> Seq Scan on heroes (cost=0.00..2.13 rows=113 width=14) (actual time=0.022..1.166 rows=113 loops=1)
29	Planning Time: 0.513 ms
30	Execution Time: 74239.906 ms

Summary

ORM generated query is absolutely identical to the original one, so it is impossible to get different performance from it.

Generally, the most expensive operations are JOIN. The execution time of initial query is 74414 ms and the execution time of the ORM generated one – 74239 ms. The difference in execution time can be considered as statistical deviation. Total cost is 154372 of every query.